

## EKS- Day#2 Task

EKS is a fully managed service which is managed by AWS.

It is highly integrated with AWS services such as EC2, EBS, ELB, EFS, and VPC

**Spot Instances:** Below yaml shows how to set up an eks cluster with a mix of spot instances & on demand instances.

Spot instances (slaves nodes) are launched only when bid price is above the spot price. We can set the bid price for spot instances while launching it.

The instances will be launched only when bid price is above the spot price. If the spot price is above bid price Spot instances will be terminated by AWS. There is a pricing associated with Spot instances.

Using eksctl we can create a single node group with mixed instances such as t2.small, t2.micro etc.

```
apiVersion: eksctl.io/v1alpha5
kind: ClusterConfig

metadata:
  name: mncluster
  region: ap-south-1

nodeGroups:
- name: ng1
  desiredCapacity: 2
  instanceType: t2.micro
  ssh:
    publicKeyName: kubeks
- name: ng2
  desiredCapacity: 1
  instanceType: t2.small
  ssh:
    publicKeyName: kubeks
- name: ng-mixed
  minSize: 2
  maxSize: 5
  instancesDistribution:
    maxPrice: 0.017
    instanceTypes: ["t3.small", "t3.medium"] # At least one instance type should be specified
    onDemandBaseCapacity: 0
    onDemandPercentageAboveBaseCapacity: 50
    spotInstancePools: 2
  ssh:
    publicKeyName: kubeks
```

## EKS Cluster creation through command line

```
Command Prompt - eksctl create cluster -f D:\EKS\D2\cluster.yml

C:\Program Files\Kubernetes\Minikube>eksctl create cluster -f D:\EKS\D2\cluster.yml
[0] eksctl version 0.21.0
[0] using region ap-south-1
[0] setting availability zones to [ap-south-1a ap-south-1c ap-south-1b]
[0] subnets for ap-south-1a - public:192.168.0.0/19 private:192.168.96.0/19
[0] subnets for ap-south-1c - public:192.168.32.0/19 private:192.168.128.0/19
[0] subnets for ap-south-1b - public:192.168.64.0/19 private:192.168.160.0/19
[0] nodegroup "ng1" will use "ami-073969767527f7306" [AmazonLinux2/1.16]
[0] using EC2 key pair "kubeks"
[0] nodegroup "ng2" will use "ami-073969767527f7306" [AmazonLinux2/1.16]
[0] using EC2 key pair "kubeks"
[0] nodegroup "ng-mixed" will use "ami-073969767527f7306" [AmazonLinux2/1.16]
[0] using EC2 key pair "kubeks"
[0] using Kubernetes version 1.16
[0] creating EKS cluster "mncluster" in "ap-south-1" region with un-managed nodes
[0] 3 nodegroups (ng-mixed, ng1, ng2) were included (based on the include/exclude rules)
[0] will create a CloudFormation stack for cluster itself and 3 nodegroup stack(s)
[0] will create a CloudFormation stack for cluster itself and 0 managed nodegroup stack(s)
[0] if you encounter any issues, check CloudFormation console or try 'eksctl utils describe-stacks --region=ap-south-1
--cluster=mncluster'
[0] CloudWatch logging will not be enabled for cluster "mncluster" in "ap-south-1"
[0] you can enable it with 'eksctl utils update-cluster-logging --region=ap-south-1 --cluster=mncluster'
[0] Kubernetes API endpoint access will use default of {publicAccess=true, privateAccess=false} for cluster "mncluster"
in "ap-south-1"
[0] 2 sequential tasks: { create cluster control plane "mncluster", 2 sequential sub-tasks: { no tasks, 3 parallel sub-
tasks: { create nodegroup "ng1", create nodegroup "ng2", create nodegroup "ng-mixed" } } }
[0] building cluster stack "eksctl-mncluster-cluster"
[0] deploying stack "eksctl-mncluster-cluster"
```

```
[0] building nodegroup stack "eksctl-mncluster-nodegroup-ng1"
[0m] building nodegroup stack "eksctl-mncluster-nodegroup-ng1"
[0m] building nodegroup stack "eksctl-mncluster-nodegroup-ng1"

[0] --nodes-min=2 was set automatically for nodegroup ng2
[0] --nodes-max=2 was set automatically for nodegroup ng1
[0] --nodes-max=1 was set automatically for nodegroup ng2
[0] --nodes-max=1 was set automatically for nodegroup ng2
[0] deploying stack "eksctl-mncluster-nodegroup-ng-mixed"
[0] deploying stack "eksctl-mncluster-nodegroup-ng1"
[0] deploying stack "eksctl-mncluster-nodegroup-ng2"
[!] retryable error (RequestError: send request failed
caused by: Post "https://cloudformation.ap-south-1.amazonaws.com/": EOF) from cloudformation/DescribeStacks - will retry after delay of 45.028752ms
[0] waiting for the control plane availability...
[!] unable to write kubeconfig, please retry with 'eksctl utils write-kubeconfig -n mncluster': unable to read existing kubeconfig file "C:\\Users\\arvind.ramugade/.k
ube/config": error loading config file "C:\\Users\\arvind.ramugade/.kube/config": read C:\\Users\\arvind.ramugade/.kube/config: The process cannot access the file because a
nother process has locked a portion of the file.
[0] no tasks
[0] all EKS cluster resources for "mncluster" have been created
[0] adding identity "arn:aws:iam::819612740240:role/eksctl-mncluster-nodegroup-ng1-NodeInstanceRole-V4N827T5Q86B" to auth ConfigMap
[0] nodegroup "ng1" has 0 node(s)
[0] waiting for at least 2 node(s) to become ready in "ng1"
[0] nodegroup "ng1" has 2 node(s)
[0] node "ip-192-168-25-50.ap-south-1.compute.internal" is ready
[0] node "ip-192-168-85-212.ap-south-1.compute.internal" is ready
[0] adding identity "arn:aws:iam::819612740240:role/eksctl-mncluster-nodegroup-ng2-NodeInstanceRole-ZFTVREN1C1MB" to auth ConfigMap
[0] nodegroup "ng2" has 0 node(s)
[0] waiting for at least 1 node(s) to become ready in "ng2"
[0] nodegroup "ng2" has 1 node(s)
[0] node "ip-192-168-74-110.ap-south-1.compute.internal" is ready
[0] adding identity "arn:aws:iam::819612740240:role/eksctl-mncluster-nodegroup-ng-mix-NodeInstanceRole-1CE72WIUHOEP2" to auth ConfigMap
[0] nodegroup "ng-mixed" has 0 node(s)
[0] waiting for at least 2 node(s) to become ready in "ng-mixed"
[0] nodegroup "ng-mixed" has 2 node(s)
[0] node "ip-192-168-14-107.ap-south-1.compute.internal" is ready
[0] node "ip-192-168-95-56.ap-south-1.compute.internal" is ready
[0] EKS cluster "mncluster" in "ap-south-1" region is ready
```

## EKS Cluster verification through command line

```
C:\Program Files\Kubernetes\Minikube>eksctl get cluster
NAME          REGION
mncluster     ap-south-1
```

```
C:\Program Files\Kubernetes\Minikube>eksctl get nodegroup --cluster mncluster
```

CLUSTER	NODEGROUP	CREATED	MIN SIZE	MAX SIZE	DESIRED CAPACITY	INSTANCE TYPE	IMAGE ID
mncluster	ng-mixed	2020-07-09T08:00:28Z	2	5	0	t3.small	ami-073969767527f7306
mncluster	ng1	2020-07-09T08:00:28Z	2	2	2	t2.micro	ami-073969767527f7306
mncluster	ng2	2020-07-09T08:00:29Z	1	1	1	t2.small	ami-073969767527f7306

```
C:\Program Files\Kubernetes\Minikube>kubectl cluster-info
Kubernetes master is running at https://7E1C806BAF6E83811E03E55720D69C41.yl4.ap-south-1.eks.amazonaws.com
CoreDNS is running at https://7E1C806BAF6E83811E03E55720D69C41.yl4.ap-south-1.eks.amazonaws.com/api/v1/namespaces/kube-system/services/kube-dns:dns/proxy
To further debug and diagnose cluster problems, use 'kubectl cluster-info dump'.
```

```
C:\Program Files\Kubernetes\Minikube>kubectl config view
apiVersion: v1
clusters:
- cluster:
  certificate-authority-data: DATA+OMITTED
  server: https://B281A5F495624F50E70A2F5B80BBBBF5.yl4.ap-south-1.eks.amazonaws.com
  name: arn:aws:eks:ap-south-1:819612740240:cluster/lwcluster
- cluster:
  certificate-authority-data: DATA+OMITTED
  server: https://7E1C806BAF6E83811E03E55720D69C41.yl4.ap-south-1.eks.amazonaws.com
  name: arn:aws:eks:ap-south-1:819612740240:cluster/mncluster
- cluster:
  certificate-authority-data: DATA+OMITTED
  server: https://F1D498ED792166C8A8DC0775B26535A2.sk1.ap-southeast-1.eks.amazonaws.com
  name: arn:aws:eks:ap-southeast-1:819612740240:cluster/far-lwcluster
- cluster:
  certificate-authority: C:\Users\arvind.ramugade\.minikube\ca.crt
  server: https://192.168.99.100:8443
  name: minikube
contexts:
- context:
  cluster: arn:aws:eks:ap-south-1:819612740240:cluster/lwcluster
  user: arn:aws:eks:ap-south-1:819612740240:cluster/lwcluster
  name: arn:aws:eks:ap-south-1:819612740240:cluster/lwcluster
- context:
  cluster: arn:aws:eks:ap-south-1:819612740240:cluster/mncluster
  namespace: mnns
  user: arn:aws:eks:ap-south-1:819612740240:cluster/mncluster
  name: arn:aws:eks:ap-south-1:819612740240:cluster/mncluster
- context:
  cluster: arn:aws:eks:ap-southeast-1:819612740240:cluster/far-lwcluster
  user: arn:aws:eks:ap-southeast-1:819612740240:cluster/far-lwcluster
  name: arn:aws:eks:ap-southeast-1:819612740240:cluster/far-lwcluster
- context:
  cluster: minikube
  user: minikube
  name: minikube
current-context: arn:aws:eks:ap-south-1:819612740240:cluster/mncluster
kind: Config
preferences: {}
users:
```

```

- name: arn:aws:eks:ap-south-1:819612740240:cluster/lwcluster
  user:
    exec:
      apiVersion: client.authentication.k8s.io/v1alpha1
      args:
        - --region
        - ap-south-1
        - eks
        - get-token
        - --cluster-name
        - lwcluster
      command: aws
      env: null
- name: arn:aws:eks:ap-south-1:819612740240:cluster/mncluster
  user:
    exec:
      apiVersion: client.authentication.k8s.io/v1alpha1
      args:
        - --region
        - ap-south-1
        - eks
        - get-token
        - --cluster-name
        - mncluster
      command: aws
      env: null
- name: arn:aws:eks:ap-southeast-1:819612740240:cluster/far-lwcluster
  user:
    exec:
      apiVersion: client.authentication.k8s.io/v1alpha1
      args:
        - --region
        - ap-southeast-1
        - eks
        - get-token
        - --cluster-name
        - far-lwcluster
      command: aws
      env: null
- name: minikube
  user:
    client-certificate: C:\Users\arvind.ramugade\.minikube\profiles\minikube\client.crt
    client-key: C:\Users\arvind.ramugade\.minikube\profiles\minikube\client.key

```

We can verify the cluster creation through AWS Console as well as shown below

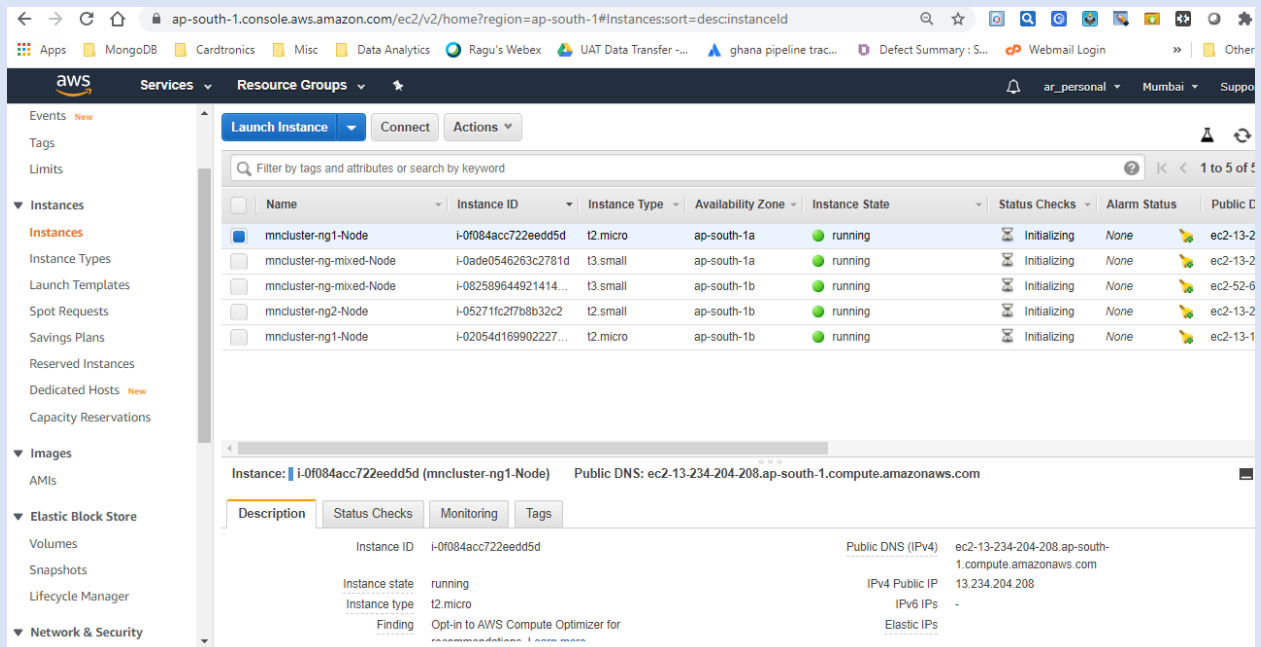
The screenshot displays the AWS Management Console interface for an Amazon EKS cluster named 'mncluster'. The left sidebar shows the navigation menu with 'Amazon EKS' and 'Clusters' selected. The main content area shows the 'Cluster configuration' section with the following details:

- Kubernetes version:** 1.16 (Status: Active)
- Platform version:** eks.1

Below the configuration section, there are tabs for 'Details', 'Compute', 'Networking', 'Logging', 'Updates', and 'Tags'. The 'Details' tab is currently selected, showing the following information:

- API server endpoint:** <https://7E1C806BAF6E83811E03E55720D69C41.y4.ap-south-1.eks.amazonaws.com>
- OpenID Connect provider URL:** <https://oidc.eks.ap-south-1.amazonaws.com/id/7E1C806BAF6E83811E03E55720D69C41>
- Cluster ARN:** arn:aws:eks:ap-south-1:819612740240:cluster/mncluster
- Creation time:** Jul 9th 2020 at 1:19 PM
- Certificate authority:** LS0tLS1CRUdJTiBDRVJUSUZJQ0FURSOtLS0tCk1JSUN5RENDQWJDZ0F3SUJBZ0lCQURBTkJn a3Foa2lHOXcwQkFRc0ZBREFWTVJNd0VRWU
- Cluster IAM Role ARN:** arn:aws:iam::819612740240:role/eksctl-mncluster-cluster-ServiceRole-M1OYWK3NMG1R

The footer of the console shows the date '© 2008 - 2020, Amazon Internet Services Private Ltd. or its affiliates. All rights reserved.' and links to 'Privacy Policy' and 'Terms of Use'.



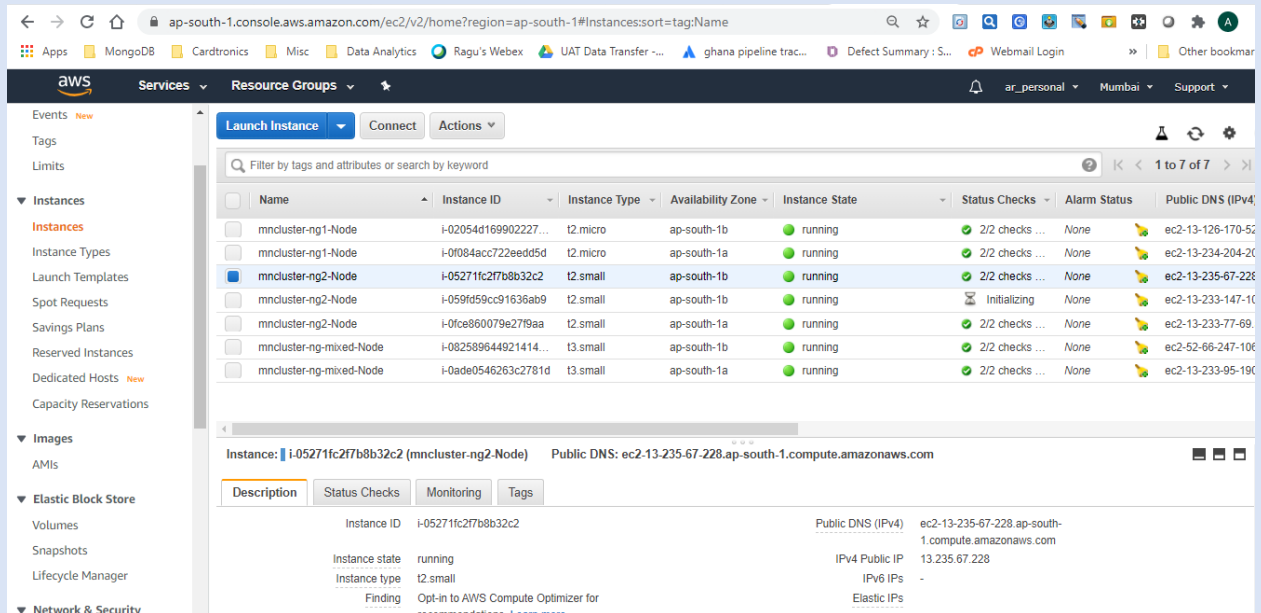
Using below command we can update kube-config file so that we can use kubectl command in EKS cluster

```
aws eks update-kubeconfig --name=mncluster
```

```
kubectl get nodes
```

We can scale our cluster for a particular node group as under:

```
C:\Program Files\Kubernetes\Minikube>eksctl scale nodegroup --cluster mncluster --name ng2 --nodes=3 --nodes-max=5
[+] scaling nodegroup stack "eksctl-mncluster-nodegroup-ng2" in cluster eksctl-mncluster-cluster
[+] scaling nodegroup, desired capacity from 1 to 3, max size from 1 to 5
```



```
C:\Program Files\Kubernetes\Minikube>eksctl get nodegroup --cluster mncluster
CLUSTER  NODEGROUP  CREATED          MIN SIZE  MAX SIZE  DESIRED CAPACITY  INSTANCE TYPE  IMAGE ID
mncluster  ng-mixed    2020-07-09T08:00:28Z  2         5         0                 t3.small       ami-073969767527f7306
mncluster  ng1         2020-07-09T08:00:28Z  2         2         2                 t2.micro       ami-073969767527f7306
mncluster  ng2         2020-07-09T08:00:29Z  1         5         3                 t2.small       ami-073969767527f7306
```

```
C:\Program Files\Kubernetes\Minikube>
```

```
C:\Program Files\Kubernetes\Minikube>kubectl create namespace mnns
namespace/mnns created

C:\Program Files\Kubernetes\Minikube>kubectl config set-context --current --namespace=mnns
Context "arn:aws:eks:ap-south-1:819612740240:cluster/mncluster" modified.
```

```
C:\Program Files\Kubernetes\Minikube>
```

By default in a container there is no connectivity for pods running on multiple nodes. If there are multiple pods in a single slave node they can communicate with each other, however they can't connect with pods in another slave.

Using CNI (also known as flannel) we can achieve this. CNI, VPC, Subnet created by Eksctl automatically once we set up them multi node cluster (1master& 2 slaves)

k8s\_coredns manages outside network connectivity.

Limit on no. of pods in a node:

Following command shows maximum no. of pods which we can launch

```
ps aux | grep kubectl
```

The limit on no. of pods which can be launched in a node varies based on instance type. e.g. for t2.micro instance it has capability of 4 NIC of which two are used for instance IP address only two NIC are available and we can run only two pods in the t2.micro.

Network interface attached by AWS CNI for pods help in inter connection with other pods on different nodes.

```
root@ip-192-168-85-212:~  
[root@ip-192-168-85-212 ~]# ps aux | grep kubelet  
root      4576   1.5  8.8 864464 89064 ?        Ssl  08:04   0:03 /usr/bin/kubelet --node-ip=192.168.85.212 --node-labels=alpha.eksctl.io/cluster-name=mncluster,alpha.eksctl.io/nodegroup-name=ngl,alpha.eksctl.io/instance-id=i-02054dl6990222776 --max-pods=4 --register-node=true --register-with-taints= --cloud-provider=aws --container-runtime=docker --network-plugin=cni --cni-bin-dir=/opt/cni/bin --cni-conf-dir=/etc/cni/net.d --pod-infra-container-image=602401143452.dkr.ecr.ap-south-1.amazonaws.com/eks/pause-amd64:3.1 --kubeconfig=/etc/eksctl/kubeconfig.yaml --config=/etc/eksctl/kubelet.yaml  
root      11112   0.0  0.0 119420   944 pts/0    S+   08:07   0:00 grep --color=auto kubelet  
[root@ip-192-168-85-212 ~]#
```

```
root@ip-192-168-85-212:~  
[root@ip-192-168-85-212 ~]# ps aux | grep kubelet  
root      4576   1.5  8.8 864464 89064 ?        Ssl  08:04   0:03 /usr/bin/kubelet --node-ip=192.168.85.212 --node-labels=alpha.eksctl.io/cluster-name=mncluster,alpha.eksctl.io/nodegroup-name=ngl,alpha.eksctl.io/instance-id=i-02054dl6990222776 --max-pods=4 --register-node=true --register-with-taints= --cloud-provider=aws --container-runtime=docker --network-plugin=cni --cni-bin-dir=/opt/cni/bin --cni-conf-dir=/etc/cni/net.d --pod-infra-container-image=602401143452.dkr.ecr.ap-south-1.amazonaws.com/eks/pause-amd64:3.1 --kubeconfig=/etc/eksctl/kubeconfig.yaml --config=/etc/eksctl/kubelet.yaml  
root      11112   0.0  0.0 119420   944 pts/0    S+   08:07   0:00 grep --color=auto kubelet  
[root@ip-192-168-85-212 ~]#
```

```
root@ip-192-168-85-212:~  
[root@ip-192-168-85-212 ~]# ifconfig -a  
eni36209f359b9: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 9001  
    inet6 fe80::6821:3eff:fe49:8bb6 prefixlen 64 scopeid 0x20<link>  
    ether 6a:21:3e:49:8b:b6 txqueuelen 0 (Ethernet)  
    RX packets 428 bytes 36451 (35.5 KiB)  
    RX errors 0 dropped 0 overruns 0 frame 0  
    TX packets 427 bytes 144354 (140.9 KiB)  
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0  
  
enibc3cd0f2e6b: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 9001  
    inet6 fe80::d032:aff:fe73:9483 prefixlen 64 scopeid 0x20<link>  
    ether d2:32:0a:73:94:83 txqueuelen 0 (Ethernet)  
    RX packets 475 bytes 40346 (39.4 KiB)  
    RX errors 0 dropped 0 overruns 0 frame 0  
    TX packets 477 bytes 158569 (154.8 KiB)  
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0  
  
eth0: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 9001  
    inet 192.168.85.212 netmask 255.255.224.0 broadcast 192.168.95.255  
    inet6 fe80::8bd:d3ff:fe05:1130 prefixlen 64 scopeid 0x20<link>  
    ether 0a:bd:d3:05:11:30 txqueuelen 1000 (Ethernet)  
    RX packets 153354 bytes 219122950 (208.9 MiB)  
    RX errors 0 dropped 0 overruns 0 frame 0  
    TX packets 9653 bytes 996491 (973.1 KiB)  
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0  
  
eth1: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 9001  
    inet 192.168.66.4 netmask 255.255.224.0 broadcast 192.168.95.255  
    inet6 fe80::8ff:39ff:fe4c:b8b4 prefixlen 64 scopeid 0x20<link>  
    ether 0a:ff:39:4c:b8:b4 txqueuelen 1000 (Ethernet)  
    RX packets 280 bytes 126361 (123.3 KiB)  
    RX errors 0 dropped 0 overruns 0 frame 0  
    TX packets 275 bytes 21708 (21.1 KiB)  
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0  
  
lo: flags=73<UP,LOOPBACK,RUNNING> mtu 65536  
    inet 127.0.0.1 netmask 255.0.0.0  
    inet6 ::1 prefixlen 128 scopeid 0x10<host>  
    loop txqueuelen 1000 (Local Loopback)  
    RX packets 1514 bytes 115587 (112.8 KiB)  
    RX errors 0 dropped 0 overruns 0 frame 0  
    TX packets 1514 bytes 115587 (112.8 KiB)  
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0
```

We can confirm that **Docker is already installed** in the instance due to EKS cluster.



```

[root@ip-192-168-85-212 ~]# systemctl status docker
● docker.service - Docker Application Container Engine
   Loaded: loaded (/usr/lib/systemd/system/docker.service; enabled; vendor preset: disabled)
   Active: active (running) since Thu 2020-07-09 08:04:27 UTC; 6min ago
     Docs: https://docs.docker.com
   Main PID: 3964 (dockerd)
    Tasks: 13
   Memory: 367.6M
   CGroup: /system.slice/docker.service
           └─3964 /usr/bin/dockerd -H fd:// --containerd=/run/containerd/containerd.sock

Jul 09 08:05:25 ip-192-168-85-212.ap-south-1.compute.internal dockerd[3964]: time="2020-07-09T08:05:25.370736969Z" level=info msg="ignoring event" module=libco...Delete
Jul 09 08:05:26 ip-192-168-85-212.ap-south-1.compute.internal dockerd[3964]: time="2020-07-09T08:05:26.374255808Z" level=info msg="ignoring event" module=libco...Delete
Jul 09 08:05:27 ip-192-168-85-212.ap-south-1.compute.internal dockerd[3964]: time="2020-07-09T08:05:27.411124168Z" level=info msg="ignoring event" module=libco...Delete
Jul 09 08:05:28 ip-192-168-85-212.ap-south-1.compute.internal dockerd[3964]: time="2020-07-09T08:05:28.414021207Z" level=info msg="ignoring event" module=libco...Delete
Jul 09 08:05:29 ip-192-168-85-212.ap-south-1.compute.internal dockerd[3964]: time="2020-07-09T08:05:29.470419680Z" level=info msg="ignoring event" module=libco...Delete
Jul 09 08:05:30 ip-192-168-85-212.ap-south-1.compute.internal dockerd[3964]: time="2020-07-09T08:05:30.489866328Z" level=info msg="ignoring event" module=libco...Delete
Jul 09 08:05:31 ip-192-168-85-212.ap-south-1.compute.internal dockerd[3964]: time="2020-07-09T08:05:31.508950762Z" level=info msg="ignoring event" module=libco...Delete
Jul 09 08:05:32 ip-192-168-85-212.ap-south-1.compute.internal dockerd[3964]: time="2020-07-09T08:05:32.523186266Z" level=info msg="ignoring event" module=libco...Delete
Jul 09 08:05:33 ip-192-168-85-212.ap-south-1.compute.internal dockerd[3964]: time="2020-07-09T08:05:33.572789446Z" level=info msg="ignoring event" module=libco...Delete
Jul 09 08:05:34 ip-192-168-85-212.ap-south-1.compute.internal dockerd[3964]: time="2020-07-09T08:05:34.610436339Z" level=info msg="ignoring event" module=libco...Delete
Hint: Some lines were ellipsized, use -l to show in full.
[root@ip-192-168-85-212 ~]#

```

```

C:\Program Files\Kubernetes\Minikube>aws eks update-kubeconfig --name=mncluster
Added new context arn:aws:eks:ap-south-1:819612740240:cluster/mncluster to C:\Users\arvind.ramugade\.kube\config

```

```

D:\EKS\D2>kubect1 create -k .
secret/mysql-pass-md7m2684d9 created
service/wordpress-mysql created
service/wordpress created
deployment.apps/wordpress-mysql created
deployment.apps/wordpress created

```

```

D:\EKS\D2>kubect1 get pods
NAME                                READY   STATUS    RESTARTS   AGE
wordpress-7b5d46f9b9-9clhq         0/1     Pending   0           3m52s
wordpress-mysql-85b46f8cd7-jj2jp   0/1     Pending   0           3m52s

```

```

D:\EKS\D2>kubect1 create -f D:\EKS\D2\create-storage.yaml
storageclass.storage.k8s.io/aws-efs created
persistentvolumeclaim/efs-wp created
persistentvolumeclaim/efs-sql created

```

```

D:\EKS\D2>kubect1 get pvc
NAME      STATUS    VOLUME   CAPACITY   ACCESS MODES   STORAGECLASS   AGE
efs-sql   Pending                                     aws-efs        27s
efs-wp    Pending                                     aws-efs        27s

```

```

D:\EKS\D2>kubect1 get pv
No resources found in mnns namespace.

```

```

D:\EKS\D2>kubect1 get pods
NAME                                READY   STATUS    RESTARTS   AGE
wordpress-7b5d46f9b9-9clhq         0/1     Pending   0           9m29s
wordpress-mysql-85b46f8cd7-jj2jp   0/1     Pending   0           9m29s

```



**EKS Supports LoadBalancer service** which automatically balances traffic on slave nodes.

Kubernetes service type is LoadBalancer & by default it uses Classic Load Balancer which are public facing.

`service.beta.kubernetes.io/aws-load-balancer-type: nlb`

Using annotations we can manage LoadBalancer configuration.

The screenshot displays two screenshots of the AWS Management Console. The top screenshot shows the 'Load balancers' page in the 'ap-south-1' region. A single Classic Load Balancer is listed with the name 'aee2fe75333324046b36b45677a53d7a'. Below the list, the 'Basic Configuration' tab is selected, showing details such as the DNS name 'aee2fe75333324046b36b45677a53d7a-ap-south-1.elb.amazonaws.com (A Record)', Type 'Classic (Migrate Now)', Scheme 'internet-facing', and Availability Zones 'subnet-0468ab363881d38a5 - ap-south-1a, subnet-077277ae0c30789f6 - ap-south-1b'. The bottom screenshot shows the 'Subnets' page in the same region. A table lists several subnets, all in an 'available' state, associated with VPC 'vpc-074b6f2c3c6aad89e'. The subnets include 'eksctl-mncluster-cluster/SubnetPrivateAPSOUTH1C', 'eksctl-mncluster-cluster/SubnetPublicAPSOUTH1B', 'eksctl-mncluster-cluster/SubnetPublicAPSOUTH1A', 'eksctl-mncluster-cluster/SubnetPrivateAPSOUTH1A', 'eksctl-mncluster-cluster/SubnetPrivateAPSOUTH1B', 'eksctl-mncluster-cluster/SubnetPublicAPSOUTH1C', 'subnet-2b5de150', and 'subnet-b14d49d9'.

Name	DNS name	State	VPC ID	Availability Zones	Type
aee2fe75333324046b36b45677a53d7a	aee2fe75333324046b36b45677a53d7a-ap-south-1.elb.amazonaws.com (A Record)	available	vpc-074b6f2c3c6aad89e	ap-south-1a, ap-south-1b	classic

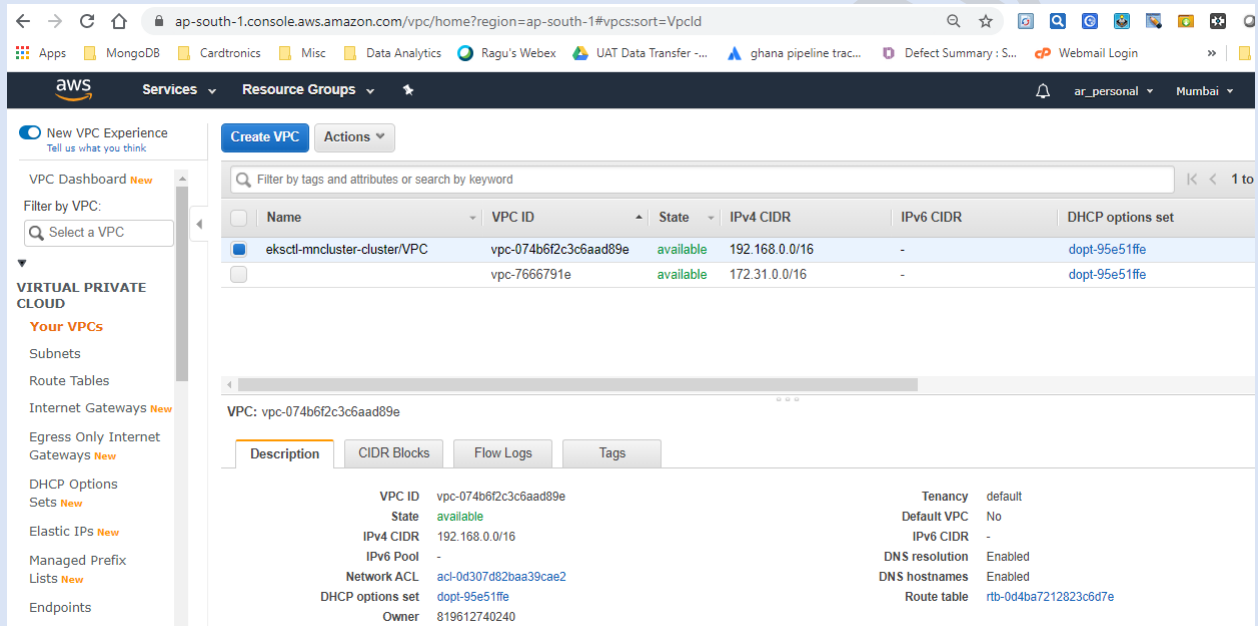
  

Name	Subnet ID	State	VPC	IPv4 CIDR
eksctl-mncluster-cluster/SubnetPrivateAPSOUTH1C	subnet-0677862a781089206	available	vpc-074b6f2c3c6aad89e	192.168.128.0/20
eksctl-mncluster-cluster/SubnetPublicAPSOUTH1B	subnet-077277ae0c30789f6	available	vpc-074b6f2c3c6aad89e	192.168.64.0/20
eksctl-mncluster-cluster/SubnetPublicAPSOUTH1A	subnet-0468ab363881d38a5	available	vpc-074b6f2c3c6aad89e	192.168.0.0/19
eksctl-mncluster-cluster/SubnetPrivateAPSOUTH1A	subnet-09c15f8da1d8af46a	available	vpc-074b6f2c3c6aad89e	192.168.96.0/20
eksctl-mncluster-cluster/SubnetPrivateAPSOUTH1B	subnet-0eeb12d0b0121319a	available	vpc-074b6f2c3c6aad89e	192.168.160.0/20
eksctl-mncluster-cluster/SubnetPublicAPSOUTH1C	subnet-09dfb0eb2b0fb1b2a	available	vpc-074b6f2c3c6aad89e	192.168.32.0/20
subnet-2b5de150	subnet-2b5de150	available	vpc-7666791e	172.31.16.0/20
subnet-b14d49d9	subnet-b14d49d9	available	vpc-7666791e	172.31.32.0/20

## EBS challenge:

With EBS we can attach only one EBS volume at a time with the EC2 instance. EBS volume can't be used to connect to EC2 instance launched in different availability zones. This is a problem in a multi node cluster set up where multiple slave nodes can't be attached to centralized storage.

This can be resolved by using **EFS service**. **EFS is a centralized NFS storage** & it spans subnets in VPC. However, the Security group should be assigned to all the subsets while creating EFS. **In EFS multiple nodes can concurrently access the storage without any performance/operational overhead.**



The screenshot displays the AWS Management Console's VPC Dashboard for the 'ap-south-1' region. The left sidebar shows the 'VIRTUAL PRIVATE CLOUD' section with 'Your VPCs' selected. The main content area shows a list of VPCs and the detailed configuration for 'vpc-074b6f2c3c6aad89e'.

Name	VPC ID	State	IPv4 CIDR	IPv6 CIDR	DHCP options set
ekscrtl-mncluster-cluster/VPC	vpc-074b6f2c3c6aad89e	available	192.168.0.0/16	-	dopt-95e51ffe
	vpc-7666791e	available	172.31.0.0/16	-	dopt-95e51ffe

VPC: vpc-074b6f2c3c6aad89e	
Description	CIDR Blocks
VPC ID	vpc-074b6f2c3c6aad89e
State	available
IPv4 CIDR	192.168.0.0/16
IPv6 Pool	-
Network ACL	acl-0d307d82baa39cae2
DHCP options set	dopt-95e51ffe
Owner	819612740240
Tenancy	default
Default VPC	No
IPv6 CIDR	-
DNS resolution	Enabled
DNS hostnames	Enabled
Route table	rtb-0d4ba7212823c6d7e



# Mount targets

VPC	Availability Zone	Subnet	IP address	Mount target ID	Network interface ID	Security groups	Mount target state
vpc-074b6f2c3c6aad89e - eksctl-mncluster-cluster/VPC	ap-south-1c	subnet-0677862a781089206 - eksctl-mncluster-cluster/SubnetPrivateAPSOUTH1C	192.168.158.107	fsmt-e132e430	eni-0a1efa54c1358d903	sg-0485a4fc4485f35bd - eksctl-mncluster-nodegroup-ng1-SG-1NWXYOVX64WNS sg-08fb1f78f78739104 - eksctl-mncluster-cluster-ClusterSharedNodeSecurityGroup-BLAWDRQ99Y22	Available
	ap-south-1b	subnet-077277ae0c30789f6 - eksctl-mncluster-cluster/SubnetPublicAPSOUTH1B	192.168.88.52	fsmt-e332e432	eni-0a4e7d272a730717a	sg-0485a4fc4485f35bd - eksctl-mncluster-nodegroup-ng1-SG-1NWXYOVX64WNS sg-08fb1f78f78739104 - eksctl-mncluster-cluster-ClusterSharedNodeSecurityGroup-BLAWDRQ99Y22	Available
	ap-south-1a	subnet-0468ab363881d38a5 - eksctl-mncluster-cluster/SubnetPublicAPSOUTH1A	192.168.3.160	fsmt-e232e433	eni-0b515ffb71fa6e4e6	sg-0485a4fc4485f35bd - eksctl-mncluster-nodegroup-ng1-SG-1NWXYOVX64WNS sg-08fb1f78f78739104 - eksctl-mncluster-cluster-ClusterSharedNodeSecurityGroup-BLAWDRQ99Y22	Available

```
D:\EKS\D2>kubect1 create -f D:\EKS\D2\create-efs-provisioner.yaml
deployment.apps/efs-provisioner created
```

```
D:\EKS\D2>kubect1 create -f D:\EKS\D2\create-rbac.yaml
clusterrolebinding.rbac.authorization.k8s.io/nfs-provisioner-role-binding created
```

```
D:\EKS\D2>kubect1 get pods
```

NAME	READY	STATUS	RESTARTS	AGE
efs-provisioner-6f54dcd88b-4bh2p	1/1	Running	0	6m14s
wordpress-7b5d46f9b9-9clhq	0/1	ContainerCreating	0	18m
wordpress-mysql-85b46f8cd7-jj2jp	0/1	ContainerCreating	0	18m

```
D:\EKS\D2>kubect1 get pvc
```

NAME	STATUS	VOLUME	CAPACITY	ACCESS MODES	STORAGECLASS	AGE
efs-sql	Bound	pvc-61787c66-27ec-41bd-b19d-e721fefb1fcc	10Gi	RWX	aws-efs	10m
efs-wp	Bound	pvc-2946c355-1a9d-4199-93a7-eb0e7e5e240a	10Gi	RWX	aws-efs	10m

```
D:\EKS\D2>kubect1 get pv
```

NAME	CAPACITY	ACCESS MODES	RECLAIM POLICY	STATUS	CLAIM	STORAGECLASS	REASON	AGE
pvc-2946c355-1a9d-4199-93a7-eb0e7e5e240a	10Gi	RWX	Delete	Bound	mnns/efs-wp	aws-efs		13m
pvc-61787c66-27ec-41bd-b19d-e721fefb1fcc	10Gi	RWX	Delete	Bound	mnns/efs-sql	aws-efs		13m

```
D:\EKS\D2>kubect1 get pods
```

NAME	READY	STATUS	RESTARTS	AGE
efs-provisioner-6f54dcd88b-4bh2p	1/1	Running	0	20m
wordpress-7b5d46f9b9-9clhq	1/1	Running	0	32m
wordpress-mysql-85b46f8cd7-jj2jp	1/1	Running	0	32m

```
D:\EKS\D2>
```

```
D:\EKS\D2>kubect1 get secret
```

NAME	TYPE	DATA	AGE
default-token-jnxdp	kubernetes.io/service-account-token	3	42m
mysql-pass-md7m2684d9	Opaque	1	34m

```
D:\EKS\D2>
```

```
D:\EKS\D2>kubect1 get sc
```

NAME	PROVISIONER	AGE
aws-efs	lw-course/aws-efs	27m
gp2 (default)	kubernetes.io/aws-ebs	69m

```
D:\EKS\D2>kubectl get all
```

NAME	READY	STATUS	RESTARTS	AGE
pod/efs-provisioner-6f54dcd88b-4bh2p	1/1	Running	0	23m
pod/wordpress-7b5d46f9b9-9clhq	1/1	Running	0	35m
pod/wordpress-mysql-85b46f8cd7-jj2jp	1/1	Running	0	35m

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
service/wordpress	LoadBalancer	10.100.24.156	aae2fe75333324046b36b45677a53d7a-2032674444.ap-south-1.elb.amazonaws.com	80:31597/TCP	35m
service/wordpress-mysql	ClusterIP	None	<none>	3306/TCP	35m

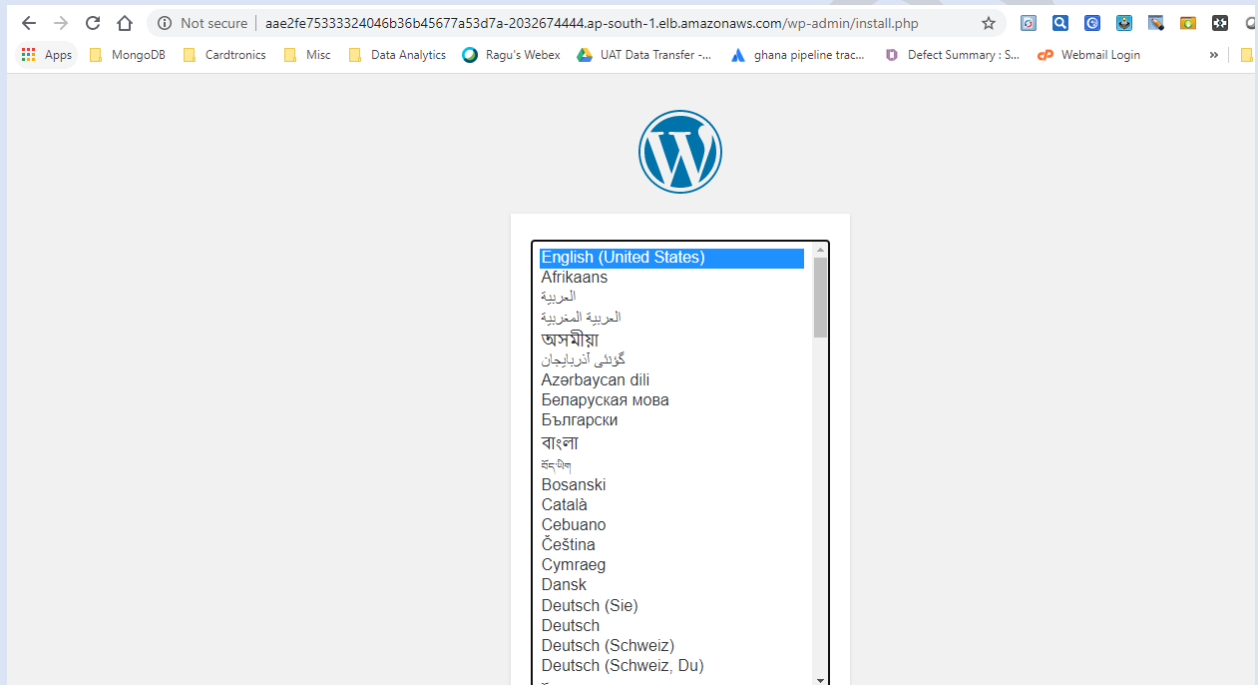
  

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
deployment.apps/efs-provisioner	1/1	1	1	23m
deployment.apps/wordpress	1/1	1	1	35m
deployment.apps/wordpress-mysql	1/1	1	1	35m

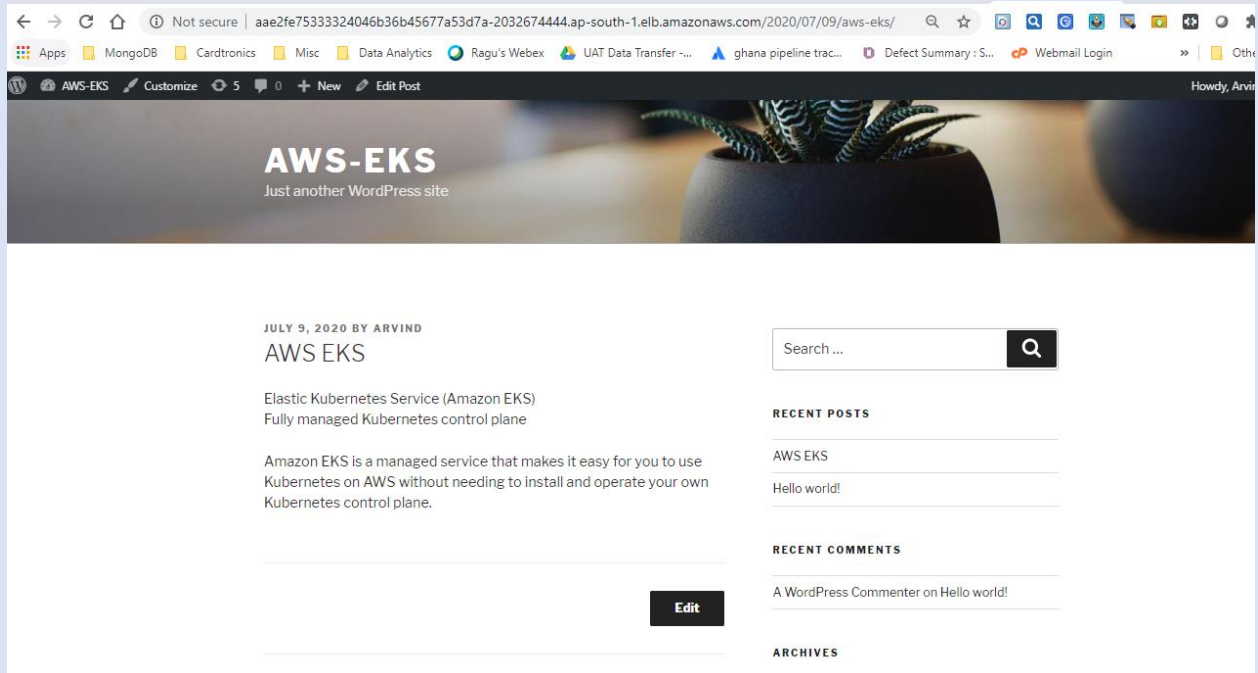
  

NAME	DESIRED	CURRENT	READY	AGE
replicaset.apps/efs-provisioner-6f54dcd88b	1	1	1	23m
replicaset.apps/wordpress-7b5d46f9b9	1	1	1	35m
replicaset.apps/wordpress-mysql-85b46f8cd7	1	1	1	35m

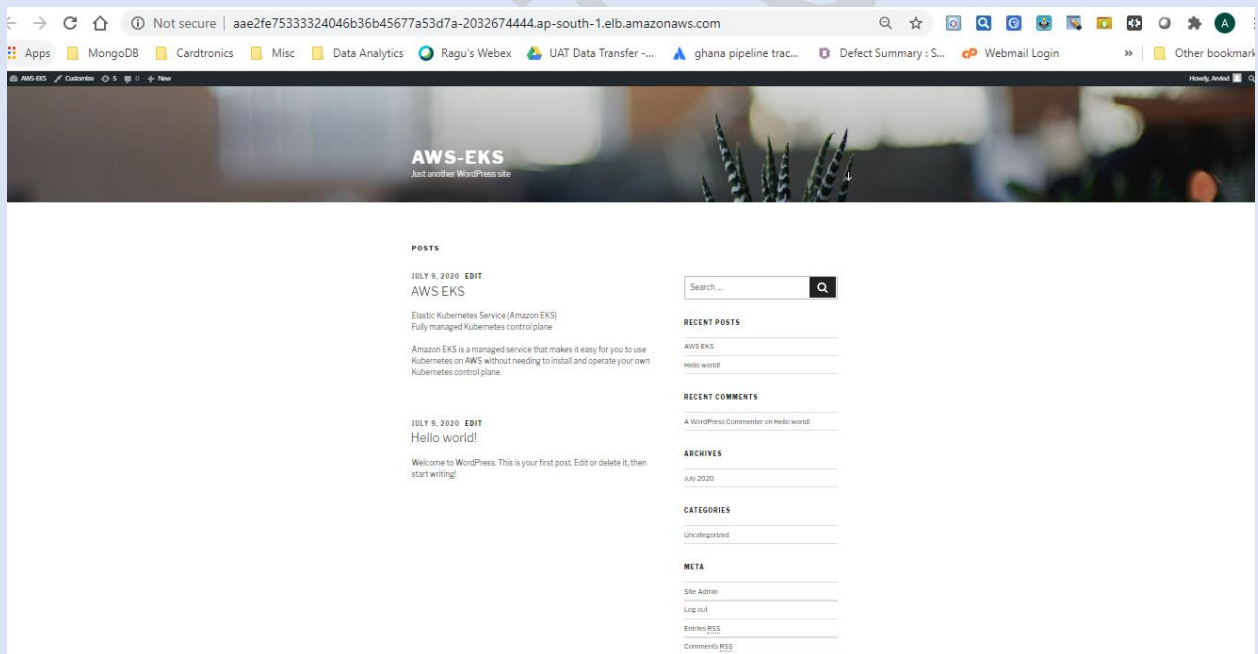
## WordPress Application



[illegible]



After clicking on LoadBalancer following page will be opened





```
D:\EKS\D2>kubectl get pods -n kube-system -o wide
NAME          READY   STATUS    RESTARTS   AGE   IP              NODE                                     NOMINATED NODE   READINESS GATES
aws-node-4pnsn 1/1     Running   0           82m   192.168.85.212  ip-192-168-85-212.ap-south-1.compute.internal <none>           <none>
aws-node-b6kr7 1/1     Running   0           70m   192.168.65.84   ip-192-168-65-84.ap-south-1.compute.internal <none>           <none>
aws-node-fgz5v 1/1     Running   0           71m   192.168.16.140  ip-192-168-16-140.ap-south-1.compute.internal <none>           <none>
aws-node-l5d9t 1/1     Running   0           80m   192.168.14.197  ip-192-168-14-197.ap-south-1.compute.internal <none>           <none>
aws-node-ng5c9 1/1     Running   0           80m   192.168.95.56   ip-192-168-95-56.ap-south-1.compute.internal <none>           <none>
aws-node-t9dn6 1/1     Running   0           81m   192.168.74.110  ip-192-168-74-110.ap-south-1.compute.internal <none>           <none>
aws-node-wxp72 1/1     Running   0           82m   192.168.25.50   ip-192-168-25-50.ap-south-1.compute.internal <none>           <none>
coredns-6856799b8d-9nfjw 1/1     Running   0           89m   192.168.82.90   ip-192-168-85-212.ap-south-1.compute.internal <none>           <none>
coredns-6856799b8d-txffd 1/1     Running   0           89m   192.168.82.102  ip-192-168-85-212.ap-south-1.compute.internal <none>           <none>
kube-proxy-6rnnj9 1/1     Running   0           80m   192.168.95.56   ip-192-168-95-56.ap-south-1.compute.internal <none>           <none>
kube-proxy-9hfvk 1/1     Running   0           71m   192.168.16.140  ip-192-168-16-140.ap-south-1.compute.internal <none>           <none>
kube-proxy-d9n5n 1/1     Running   0           82m   192.168.85.212  ip-192-168-85-212.ap-south-1.compute.internal <none>           <none>
kube-proxy-mm8rt 1/1     Running   0           70m   192.168.65.84   ip-192-168-65-84.ap-south-1.compute.internal <none>           <none>
kube-proxy-qjwsq 1/1     Running   0           81m   192.168.74.110  ip-192-168-74-110.ap-south-1.compute.internal <none>           <none>
kube-proxy-tv6vv 1/1     Running   0           82m   192.168.25.50   ip-192-168-25-50.ap-south-1.compute.internal <none>           <none>
kube-proxy-wncw9 1/1     Running   0           80m   192.168.14.197  ip-192-168-14-197.ap-south-1.compute.internal <none>           <none>
```

## Helm

Helm is a client side program that provides the k8s software packages where we can launch the whole application in the kubernetes cluster.

## Tiller

Tiller is a server side program to help the helm to set up the whole infrastructure.

## Initializing Helm

```
D:\EKS\D2>helm init
Creating C:\Users\arvind.ramugade\.helm
Creating C:\Users\arvind.ramugade\.helm\repository
Creating C:\Users\arvind.ramugade\.helm\repository\cache
Creating C:\Users\arvind.ramugade\.helm\repository\local
Creating C:\Users\arvind.ramugade\.helm\plugins
Creating C:\Users\arvind.ramugade\.helm\starters
Creating C:\Users\arvind.ramugade\.helm\cache\archive
Creating C:\Users\arvind.ramugade\.helm\repository\repositories.yaml
Adding stable repo with URL: https://kubernetes-charts.storage.googleapis.com
Adding local repo with URL: http://127.0.0.1:8879/charts
$HELM_HOME has been configured at C:\Users\arvind.ramugade\.helm.

Tiller (the Helm server-side component) has been installed into your Kubernetes Cluster.

Please note: by default, Tiller is deployed with an insecure 'allow unauthenticated users' policy.
To prevent this, run `helm init` with the --tiller-tls-verify flag.
For more information on securing your installation see: https://v2.helm.sh/docs/securing_installation/
```

```
D:\EKS\D2>helm repo list
NAME    URL
stable  https://kubernetes-charts.storage.googleapis.com
local   http://127.0.0.1:8879/charts
```

## Creating Tiller Service

```
Run 'kubectl --help' for usage.

D:\EKS\D2>
D:\EKS\D2>kubectl -n kube-system create service account tiller
error: unknown command "account tiller"
See 'kubectl create service -h' for help and examples

D:\EKS\D2>kubectl -n kube-system create serviceaccount tiller
serviceaccount/tiller created

D:\EKS\D2>kubectl create clusterrolebinding tiller --clusterrole cluster-admin --serviceaccount=kube-system:tiller
clusterrolebinding.rbac.authorization.k8s.io/tiller created
```

```
D:\EKS\D2>helm init --service-account tiller --upgrade  
$HELM_HOME has been configured at C:\Users\arvind.ramugade\.helm.
```

```
Tiller (the Helm server-side component) has been updated to gcr.io/kubernetes-helm/tiller:v2.16.9 .
```

DO NOT COPY

## Installation of Prometheus using helm

DO NOT COPY

```

D:\EKS\02>helm install stable/prometheus --namespace prometheus --set alertmanager.persistentVolume.storageClass="gp2" --set server.persistentVolume.storageClass="gp2"
NAME:      sullen-manatee
LAST DEPLOYED: Thu Jul 9 15:14:12 2020
NAMESPACE: prometheus
STATUS:    DEPLOYED

RESOURCES:
==> v1/ConfigMap
NAME                                DATA  AGE
sullen-manatee-prometheus-alertmanager  1      1s
sullen-manatee-prometheus-server        5      1s

==> v1/DaemonSet
NAME                                DESIRED  CURRENT  READY  UP-TO-DATE  AVAILABLE  NODE SELECTOR  AGE
sullen-manatee-prometheus-node-exporter  7         7         0        7           0          <none>         1s

==> v1/Deployment
NAME                                READY  UP-TO-DATE  AVAILABLE  AGE
sullen-manatee-kube-state-metrics    0/1    1           0          1s
sullen-manatee-prometheus-alertmanager 0/1    1           0          1s
sullen-manatee-prometheus-pushgateway 0/1    1           0          1s
sullen-manatee-prometheus-server      0/1    1           0          1s

==> v1/PersistentVolumeClaim
NAME                                STATUS  VOLUME  CAPACITY  ACCESS MODES  STORAGECLASS  AGE
sullen-manatee-prometheus-alertmanager  Pending  gp2     1s
sullen-manatee-prometheus-server        Pending  gp2     1s

==> v1/Pod(related)
NAME                                READY  STATUS             RESTARTS  AGE
sullen-manatee-kube-state-metrics-5f6d4997d7-qgsmv  0/1    ContainerCreating   0          1s
sullen-manatee-prometheus-alertmanager-88f6655cf-njw9  0/2    Pending             0          1s
sullen-manatee-prometheus-node-exporter-4mst5        0/1    ContainerCreating   0          1s
sullen-manatee-prometheus-node-exporter-4rdhp        0/1    ContainerCreating   0          1s
sullen-manatee-prometheus-node-exporter-cpgrm        0/1    ContainerCreating   0          1s
sullen-manatee-prometheus-node-exporter-gcdtj        0/1    ContainerCreating   0          1s
sullen-manatee-prometheus-node-exporter-jdnmc        0/1    ContainerCreating   0          1s
sullen-manatee-prometheus-node-exporter-n6n4d        0/1    Pending             0          1s
sullen-manatee-prometheus-node-exporter-whc5f        0/1    ContainerCreating   0          1s
sullen-manatee-prometheus-pushgateway-97679b887-jm5d6 0/1    ContainerCreating   0          1s
sullen-manatee-prometheus-server-64d5bdbc4-wmkfg     0/2    Pending             0          1s

sullen-manatee-prometheus-node-exporter-whc5f        0/1    ContainerCreating   0          1s
sullen-manatee-prometheus-pushgateway-97679b887-jm5d6 0/1    ContainerCreating   0          1s
sullen-manatee-prometheus-server-64d5bdbc4-wmkfg     0/2    Pending             0          1s

==> v1/Service
NAME                                TYPE        CLUSTER-IP      EXTERNAL-IP  PORT(S)  AGE
sullen-manatee-kube-state-metrics  ClusterIP   10.100.67.81    <none>       8080/TCP  1s
sullen-manatee-prometheus-alertmanager  ClusterIP   10.100.147.205  <none>       80/TCP    1s
sullen-manatee-prometheus-node-exporter  ClusterIP   None            <none>       9100/TCP  1s
sullen-manatee-prometheus-pushgateway    ClusterIP   10.100.98.165   <none>       9091/TCP  1s
sullen-manatee-prometheus-server        ClusterIP   10.100.53.51    <none>       80/TCP    1s

==> v1/ServiceAccount
NAME                                SECRETS  AGE
sullen-manatee-kube-state-metrics    1         1s
sullen-manatee-prometheus-alertmanager 1         1s
sullen-manatee-prometheus-node-exporter 1         1s
sullen-manatee-prometheus-pushgateway  1         1s
sullen-manatee-prometheus-server       1         1s

==> v1beta1/ClusterRole
NAME                                AGE
sullen-manatee-prometheus-server      1s
sullen-manatee-kube-state-metrics      1s
sullen-manatee-prometheus-alertmanager 1s
sullen-manatee-prometheus-pushgateway  1s

==> v1beta1/ClusterRoleBinding
NAME                                AGE
sullen-manatee-prometheus-server      1s
sullen-manatee-kube-state-metrics      1s
sullen-manatee-prometheus-alertmanager 1s
sullen-manatee-prometheus-pushgateway  1s

NOTES:
The Prometheus server can be accessed via port 80 on the following DNS name from within your cluster:
sullen-manatee-prometheus-server.prometheus.svc.cluster.local

Get the Prometheus server URL by running these commands in the same shell:
export POD_NAME=$(kubectl get pods --namespace prometheus -l "app=prometheus,component=server" -o jsonpath="{.items[0].metadata.name}")
kubectl --namespace prometheus port-forward $POD_NAME 9090

```

```
Get the Alertmanager URL by running these commands in the same shell:
export POD_NAME=$(kubectl get pods --namespace prometheus -l "app=prometheus,component=alertmanager" -o jsonpath="{.items[0].metadata.name}")
kubectl --namespace prometheus port-forward $POD_NAME 9093

#####
##### WARNING: Pod Security Policy has been moved to a global property. #####
##### use .Values.podSecurityPolicy.enabled with pod-based #####
##### annotations #####
##### (e.g. .Values.nodeExporter.podSecurityPolicy.annotations) #####
#####
```

```
Get the PushGateway URL by running these commands in the same shell:
export POD_NAME=$(kubectl get pods --namespace prometheus -l "app=prometheus,component=pushgateway" -o jsonpath="{.items[0].metadata.name}")
kubectl --namespace prometheus port-forward $POD_NAME 9091
```

```
D:\EKS\D2>kubectl get pods -n prometheus
```

NAME	READY	STATUS	RESTARTS	AGE
sullen-manatee-kube-state-metrics-5f6d4997d7-qgsmv	1/1	Running	0	2m33s
sullen-manatee-prometheus-alertmanager-88f6655cf-njwj9	2/2	Running	0	2m33s
sullen-manatee-prometheus-node-exporter-4mst5	1/1	Running	0	2m33s
sullen-manatee-prometheus-node-exporter-4rdhp	1/1	Running	0	2m33s
sullen-manatee-prometheus-node-exporter-cpgrm	1/1	Running	0	2m33s
sullen-manatee-prometheus-node-exporter-gcdtj	1/1	Running	0	2m33s
sullen-manatee-prometheus-node-exporter-jdnclm	1/1	Running	0	2m33s
sullen-manatee-prometheus-node-exporter-n6n4d	0/1	Pending	0	2m33s
sullen-manatee-prometheus-node-exporter-whc5f	1/1	Running	0	2m33s
sullen-manatee-prometheus-pushgateway-97679b887-jmds6	1/1	Running	0	2m33s
sullen-manatee-prometheus-server-64d5bdbdc4-wmkfg	2/2	Running	0	2m33s

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
sullen-manatee-kube-state-metrics	ClusterIP	10.100.67.81	<none>	8080/TCP	3m45s
sullen-manatee-prometheus-alertmanager	ClusterIP	10.100.147.205	<none>	80/TCP	3m45s
sullen-manatee-prometheus-node-exporter	ClusterIP	None	<none>	9100/TCP	3m45s
sullen-manatee-prometheus-pushgateway	ClusterIP	10.100.98.165	<none>	9091/TCP	3m45s
sullen-manatee-prometheus-server	ClusterIP	10.100.53.51	<none>	80/TCP	3m45s

← → ↻ 📑

🔍 127.0.0.1:8888/graph?g0.range\_input=1h&g0.expr=&g0.tab=0 ⭐

🗂️ 🏠 🖨️ 🔌 🔄 🧩 🛡️ 👤 A ⋮

Apps MongoDB Cardtronics Misc Data Analytics Ragu's Webex UAT Data Transfer ~... ghana pipeline trac... Defect Summary : S... Webmail Login » Other bookmarks

Prometheus Alerts Graph Status ▾ Help

☐ Enable query history Try experimental React UI

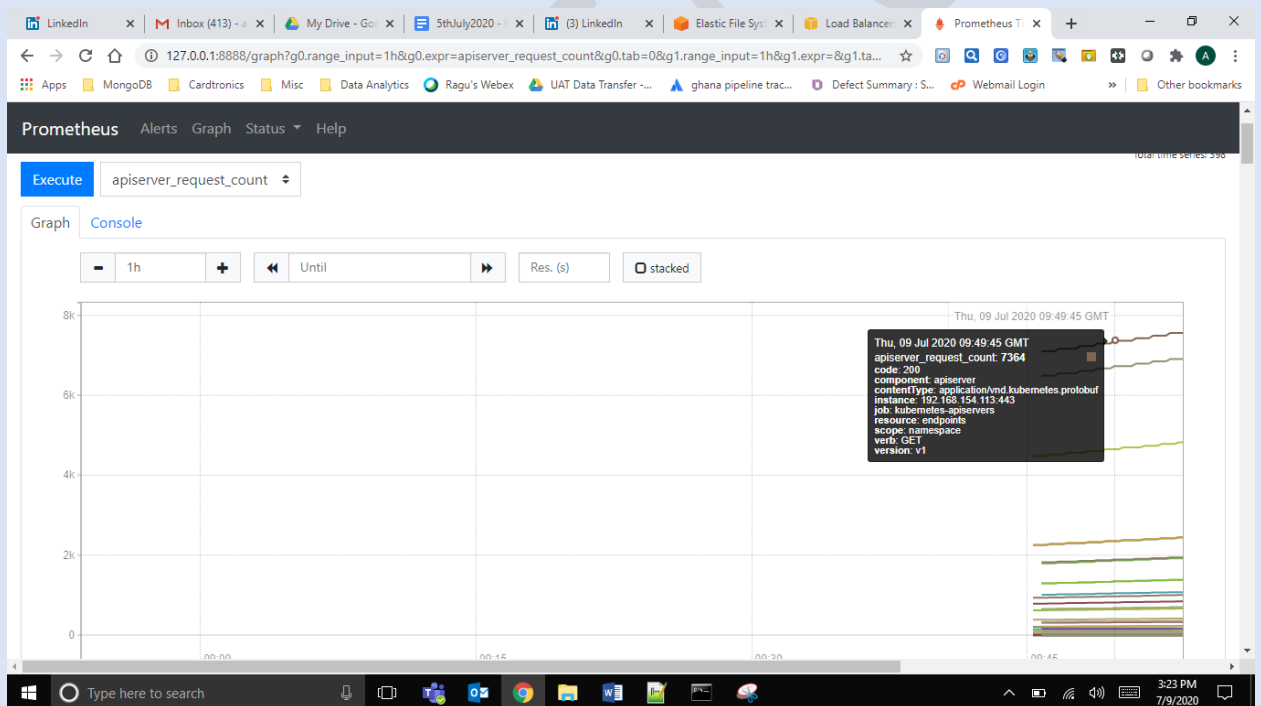
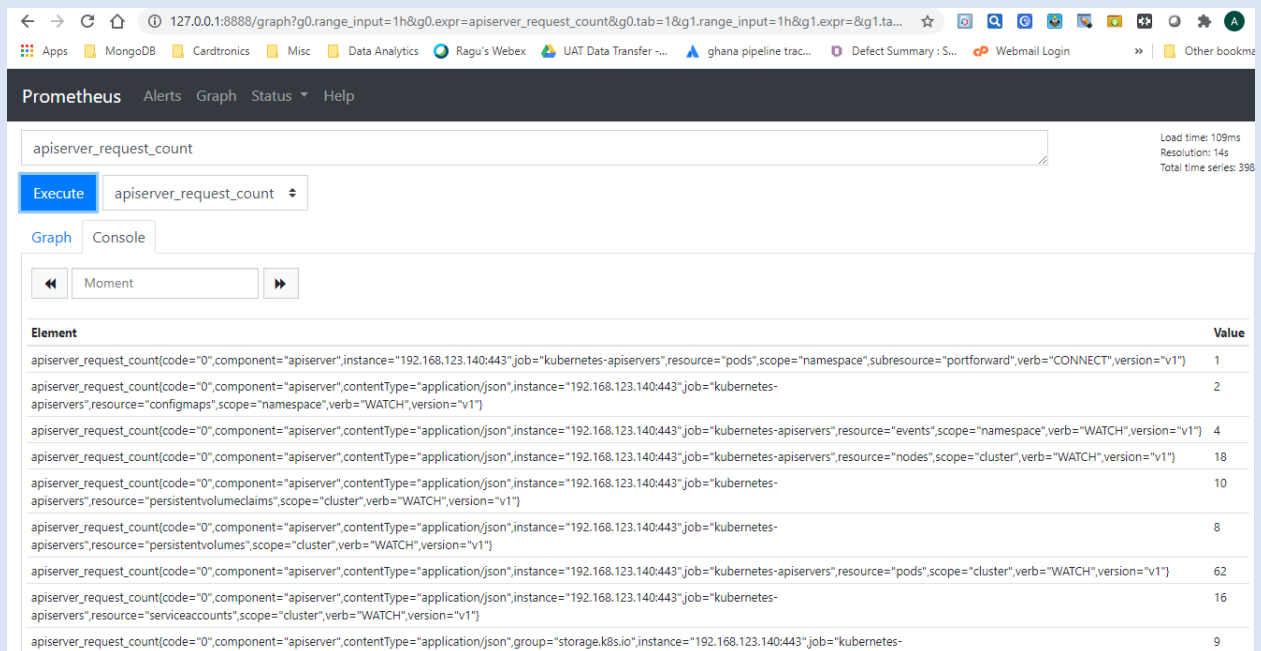
Expression (press Shift+Enter for newlines)

Execute - insert metric at cursor - ▾

Graph Console

- 1h + ◀ Until ▶ Res. (s) ☐ stacked

Add Graph Remove Graph



Prometheus Alerts Graph Status Help					
Targets					
All Unhealthy					
kubernetes-apisservers (2/2 up) show less					
Endpoint	State	Labels	Last Scrape	Scrape Duration	Error
https://192.168.123.140:443/metrics	UP	instance="192.168.123.140:443" job="kubernetes-apisservers"	50.718s ago	58.7ms	
https://192.168.154.113:443/metrics	UP	instance="192.168.154.113:443" job="kubernetes-apisservers"	18.371s ago	97.63ms	
kubernetes-nodes (7/7 up) show less					
Endpoint	State	Labels	Last Scrape	Scrape Duration	Error
https://kubernetes.default.svc:443/api/v1/nodes/ip-192-168-14-197.a p-south-1.compute.internal/proxy/ metrics	UP	alpha_eksctl_io_cluster_name="mncluster" alpha_eksctl_io_instance_id="i-0ade0546263c2781d" alpha_eksctl_io_nodegroup_name="ng-mixed" beta_kubernetes_io_arch="amd64" kubernetes_io_arch="amd64"	33.91s ago	12.19ms	
prometheus (1/1 up) show less					
Endpoint	State	Labels	Last Scrape	Scrape Duration	Error
http://localhost:9090/metrics	UP	instance="localhost:9090" job="prometheus"	30.057s ago	6.273ms	
prometheus-pushgateway (1/1 up) show less					
Endpoint	State	Labels	Last Scrape	Scrape Duration	Error
http://sullen-manatee-prometheus- pushgateway.prometheus.svc:9091/ metrics	UP	instance="sullen-manatee-prometheus-pushgateway.prometheus.svc:9091" job="prometheus-pushgateway"	36.592s ago	2.861ms	

Install Grafana.



```
D:\EKS\02>helm install stable/grafana --namespace grafana --set persistent.storageClassName="gp2" --set adminPassword=bingos --set service.type=LoadBalancer
NAME:      smelly-hummingbird
LAST DEPLOYED: Thu Jul 9 15:34:32 2020
NAMESPACE: grafana
STATUS:    DEPLOYED
```

#### RESOURCES:

```
==> v1/ClusterRole
NAME
smelly-hummingbird-grafana-clusterrole    AGE
0s

==> v1/ClusterRoleBinding
NAME
smelly-hummingbird-grafana-clusterrolebinding    AGE
0s

==> v1/ConfigMap
NAME      DATA    AGE
smelly-hummingbird-grafana    1    0s
smelly-hummingbird-grafana-test    1    0s

==> v1/Deployment
NAME      READY    UP-TO-DATE    AVAILABLE    AGE
smelly-hummingbird-grafana    0/1    1    0    0s

==> v1/Pod(related)
NAME      READY    STATUS    RESTARTS    AGE
smelly-hummingbird-grafana-697cbdf855-v69xx    0/1    ContainerCreating    0    0s

==> v1/Role
NAME      AGE
smelly-hummingbird-grafana-test    0s

==> v1/RoleBinding
NAME      AGE
smelly-hummingbird-grafana-test    0s

==> v1/Secret
NAME      TYPE    DATA    AGE
smelly-hummingbird-grafana    Opaque    3    0s

==> v1/Service
NAME      TYPE    CLUSTER-IP    EXTERNAL-IP    PORT(S)    AGE
smelly-hummingbird-grafana    LoadBalancer    10.100.158.200    <pending>    80:31570/TCP    0s
```

```
==> v1/ServiceAccount
NAME      SECRETS    AGE
smelly-hummingbird-grafana    1    0s
smelly-hummingbird-grafana-test    1    0s

==> v1beta1/PodSecurityPolicy
NAME      PRIV    CAPS    SELINUX    RUNASUSER    FSGROUP    SUPGROUP    READONLYROOTFS    VOLUMES
smelly-hummingbird-grafana    false    RunAsAny    RunAsAny    RunAsAny    RunAsAny    false    configMap,emptyDir,projected,secret,downwardAPI,persistentVolumeClaim
smelly-hummingbird-grafana-test    false    RunAsAny    RunAsAny    RunAsAny    RunAsAny    false    configMap,downwardAPI,emptyDir,projected,secret

==> v1beta1/Role
NAME      AGE
smelly-hummingbird-grafana    0s

==> v1beta1/RoleBinding
NAME      AGE
smelly-hummingbird-grafana    0s
```

#### NOTES:

1. Get your 'admin' user password by running:

```
kubectl get secret --namespace grafana smelly-hummingbird-grafana -o jsonpath="{.data.admin-password}" | base64 --decode ; echo
```

2. The Grafana server can be accessed via port 80 on the following DNS name from within your cluster:

```
smelly-hummingbird-grafana.grafana.svc.cluster.local
```

Get the Grafana URL to visit by running these commands in the same shell:

NOTE: It may take a few minutes for the LoadBalancer IP to be available.

You can watch the status of by running 'kubectl get svc --namespace grafana -w smelly-hummingbird-grafana'

```
export SERVICE_IP=$(kubectl get svc --namespace grafana smelly-hummingbird-grafana -o jsonpath='{.status.loadBalancer.ingress[0].ip}')
http://$SERVICE_IP:80
```

3. Login with the password from step 1 and the username: admin

```
#####
##### WARNING: Persistence is disabled!!! You will lose your data when #####
##### the Grafana pod is terminated. #####
#####
```

```
D:\EKS\02>kubectl get all -n grafana
```

```
NAME      READY    STATUS    RESTARTS    AGE
pod/smelly-hummingbird-grafana-697cbdf855-v69xx    1/1    Running    0    5m36s


NAME      TYPE    CLUSTER-IP    EXTERNAL-IP    PORT(S)    AGE
service/smelly-hummingbird-grafana    LoadBalancer    10.100.158.200    ad5490ec323d24afea35d9fd1f39786f-247102088.ap-south-1.elb.amazonaws.com    80:31570/TCP    5m36s

NAME      READY    UP-TO-DATE    AVAILABLE    AGE
deployment.apps/smelly-hummingbird-grafana    1/1    1    1    5m36s

NAME      DESIRED    CURRENT    READY    AGE
replicaset.apps/smelly-hummingbird-grafana-697cbdf855    1    1    1    5m36s
```

← → ↻ ⌂ ⚠ Not secure | ad5490ec323d24afea35d9fd1f39786f-247102088.ap-south-1.elb.amazonaws.com/login

Apps MongoDB Cardtronics Misc Data Analytics Ragu's Webex UAT Data Transfer ... ghana pipeline trac... Defect Summary: S... Webmail Login » Other bool



## Welcome to Grafana

Democratising data

Email or username




Password

[Log in](#)

[Forgot your password?](#)

← → ↻ ⌂ ⓘ Not secure | ad5490ec323d24afea35d9fd1f39786f-247102088.ap-south-1.elb.amazonaws.com/?orgId=1

Apps MongoDB Cardtronics Misc Data Analytics Ragu's Webex UAT Data Transfer ... ghana pipeline trac... Defect Summary: S... Webmail Login » Other bookmark

**Home**

## Welcome to Grafana

Need help? [Documentation](#) [Tutorials](#) [Community](#) [Public Slack](#)


**Basic**

The steps below will guide you to quickly finish setting up your Grafana installation.

**TUTORIAL**  
**DATA SOURCE AND DASHBOARDS**


### Grafana fundamentals

Set up and understand Grafana if you have no prior experience. This tutorial guides you through the entire process and covers the "Data source" and "Dashboards" steps to the right.



**DATA SOURCES**


### Add your first data source



[Learn how in the docs](#)


**DASHBOARDS**

### Create your first dashboard



[Learn how in the docs](#)

[Remove this panel](#)



Starred dashboards

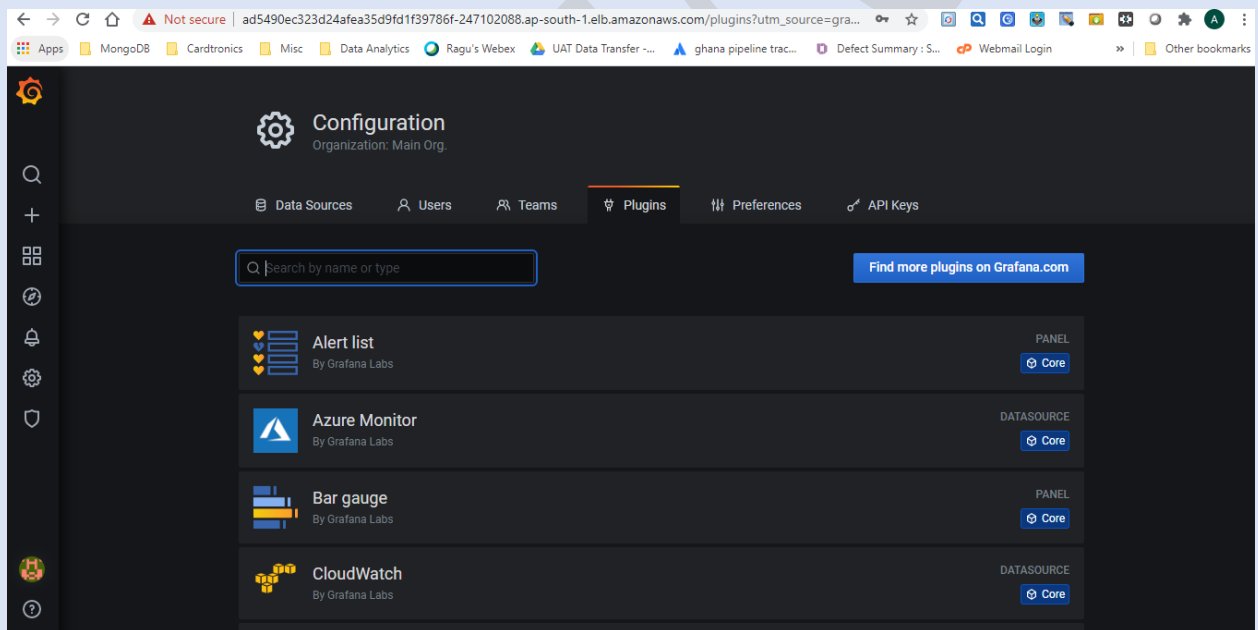
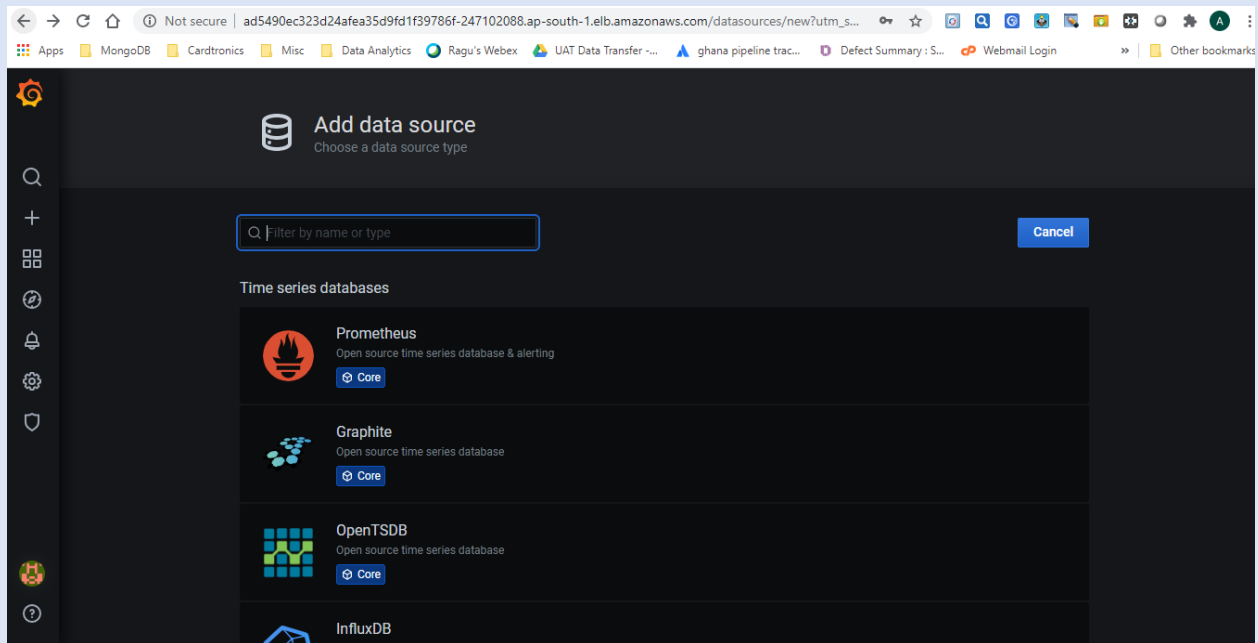
**Dashboards**

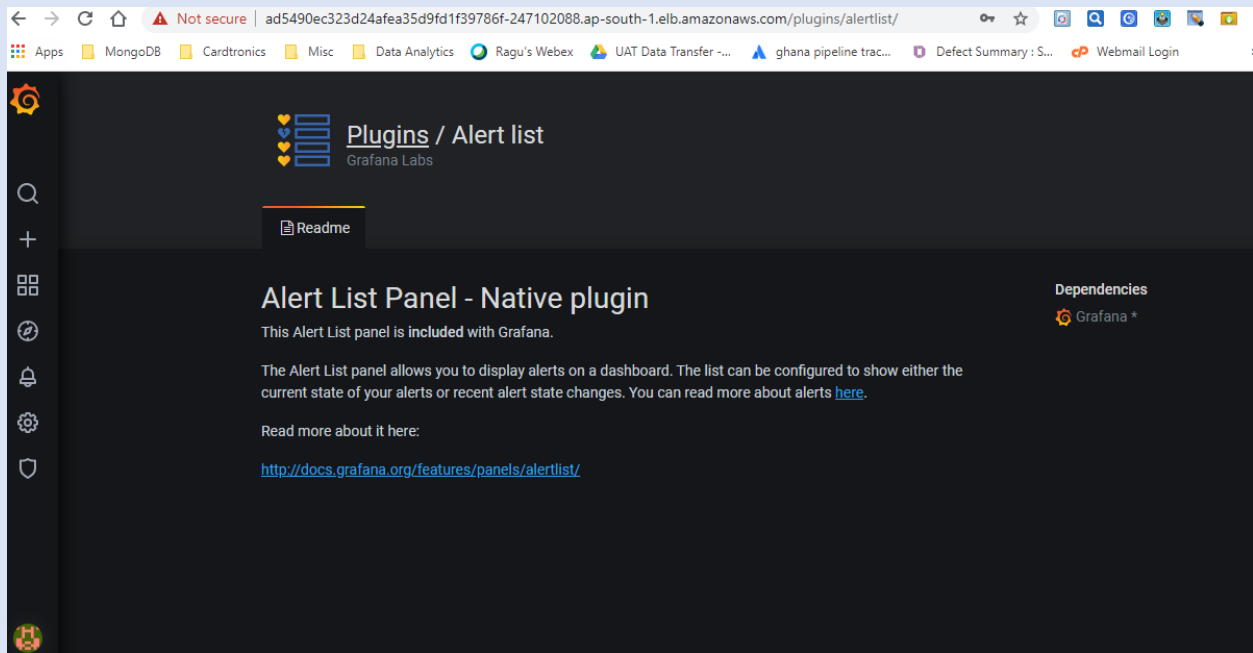
**Latest from the blog**

Grafana and NGINX are partnering to give the open source community a

Jul 08

Select Prometheus Data Source



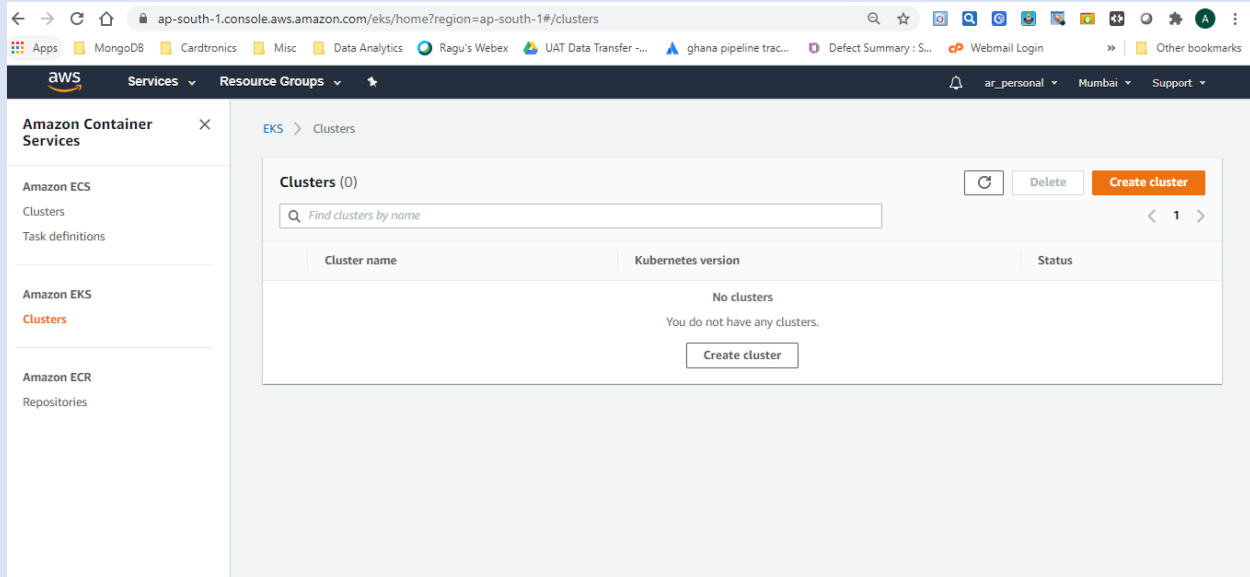


To delete entire EKS Cluster using below command

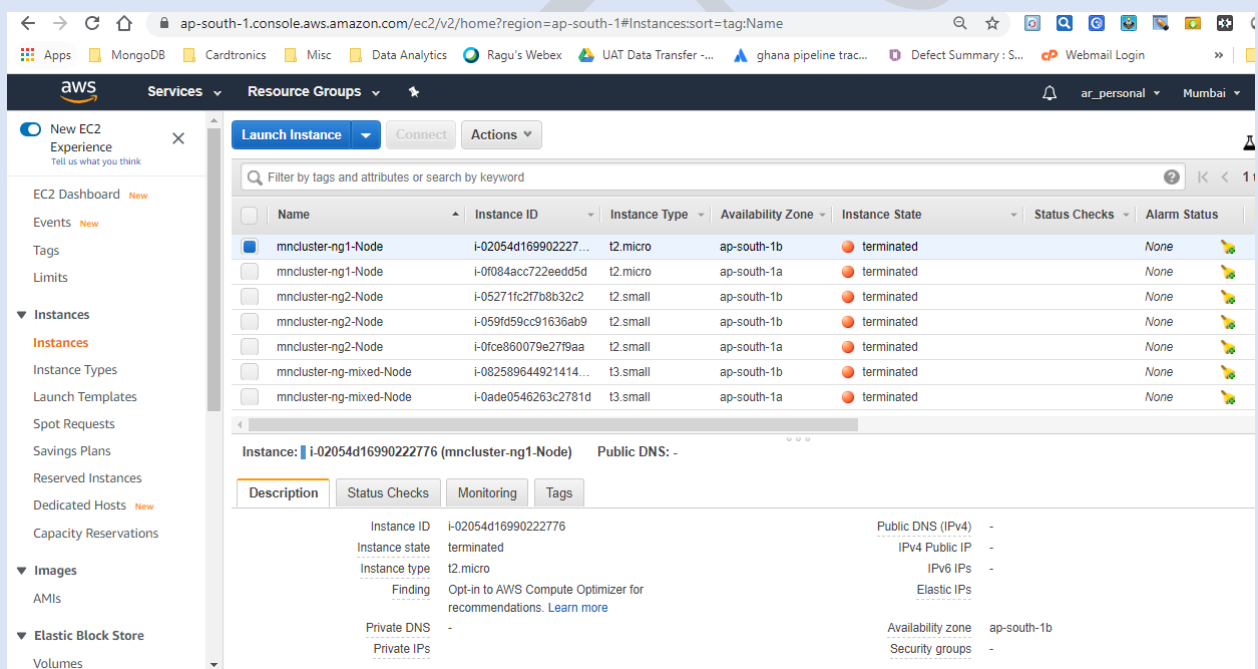
```
D:\EKS\D2>eksctl delete cluster -f cluster.yml
[0] eksctl version 0.21.0
[0] using region ap-south-1
[0] deleting EKS cluster "mncluster"
[0] either account is not authorized to use Fargate or region ap-south-1 is not supported. Ignoring error
[0] cleaning up LoadBalancer services
[0] 2 sequential tasks: { 3 parallel sub-tasks: { delete nodegroup "ng2", delete nodegroup "ng1", delete nodegroup "ng-mixed" }, delete cluster control plane "mncluster" [async] }
[0] will delete stack "eksctl-mncluster-nodegroup-ng2"
[0] waiting for stack "eksctl-mncluster-nodegroup-ng2" to get deleted
[0] will delete stack "eksctl-mncluster-nodegroup-ng-mixed"
[0] waiting for stack "eksctl-mncluster-nodegroup-ng1"
to get deleted
[0] stack "eksctl-mncluster-nodegroup-ng1"
[0] waiting for stack "eksctl-mncluster-nodegroup-ng1" to get deleted
[0] will delete stack "eksctl-mncluster-cluster"
[0] all cluster resources were deleted
```

```
D:\EKS\D2>
D:\EKS\D2>eksctl get cluster
No clusters found
```

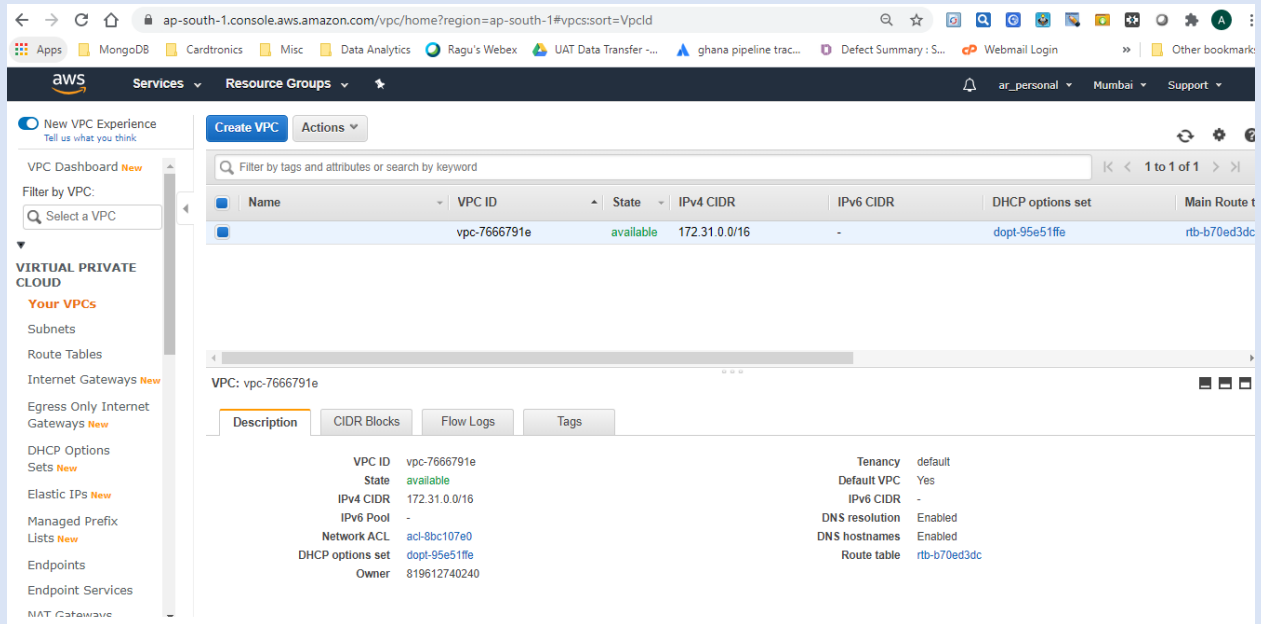
Verify deletion of cluster in AWS Console



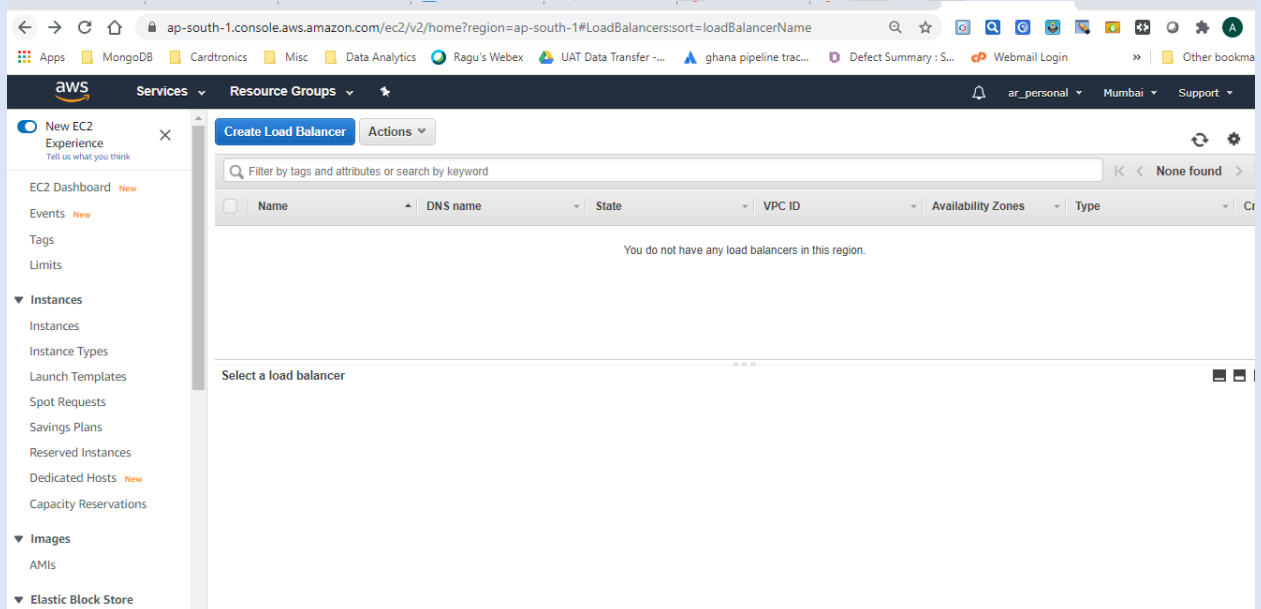
Verify that all EC2 instances are terminated.



Verify that the VPC which was created by EKS Cluster also gets removed  
Below is the default VPC which is created by AWS in a region.



Load Balancer also gets removed as part of cluster deletion.



We can remove EFS volume finally which is not required.

```
C:\Program Files\Kubernetes\Minikube>eksctl get cluster
No clusters found
```

## AWS Fargate

Fargate is a server-less architecture & it integrates with EKS.

It's a black box which dynamically manages slave nodes without worrying about capacity such as CPU, RAM etc.

It launches slaves at run time & manages internally. We can't see those slaves in the AWS console though.

```
C:\Program Files\Kubernetes\Minikube>eksctl create cluster -f D:\EKS\D2\cluster.yml
[+] eksctl version 0.21.0
[+] using region ap-southeast-1
[+] setting availability zones to [ap-southeast-1b ap-southeast-1c ap-southeast-1a]
[+] subnets for ap-southeast-1b - public:192.168.0.0/19 private:192.168.96.0/19
[+] subnets for ap-southeast-1c - public:192.168.32.0/19 private:192.168.128.0/19
[+] subnets for ap-southeast-1a - public:192.168.64.0/19 private:192.168.160.0/19
[+] using Kubernetes version 1.16
[+] creating EKS cluster "far-lwcluster" in "ap-southeast-1" region with Fargate profile
[+] will create a CloudFormation stack for cluster itself and 0 nodegroup stack(s)
[+] will create a CloudFormation stack for cluster itself and 0 managed nodegroup stack(s)
[+] if you encounter any issues, check CloudFormation console or try 'eksctl utils describe-stacks --region=ap-southeast-1 --cluster=far-lwcluster'
[+] CloudWatch logging will not be enabled for cluster "far-lwcluster" in "ap-southeast-1"
[+] you can enable it with 'eksctl utils update-cluster-logging --region=ap-southeast-1 --cluster=far-lwcluster'
[+] Kubernetes API endpoint access will use default of {publicAccess=true, privateAccess=false} for cluster "far-lwcluster" in "ap-southeast-1"
[+] 2 sequential tasks: { create cluster control plane "far-lwcluster", no tasks }
[+] building cluster stack "eksctl-far-lwcluster-cluster"
[+] deploying stack "eksctl-far-lwcluster-cluster"
[+] waiting for the control plane availability...
[!] unable to write kubeconfig , please retry with 'eksctl utils write-kubeconfig -n far-lwcluster': unable to read existing kubeconfig file "C:\Users\arvind.ramugade/.kube/config": error loading config file "C:\Users\arvind.ramugade/.kube/config": read C:\Users\arvind.ramugade/.kube/config: The process cannot access the file because another process has locked a portion of the file.
[+] no tasks
[+] all EKS cluster resources for "far-lwcluster" have been created
[+] creating Fargate profile "fargate-default" on EKS cluster "far-lwcluster"
[+] created Fargate profile "fargate-default" on EKS cluster "far-lwcluster"
[+] "coredns" is now schedulable onto Fargate
[+] "coredns" is now scheduled onto Fargate
[+] "coredns" pods are now scheduled onto Fargate
[+] EKS cluster "far-lwcluster" in "ap-southeast-1" region is ready

C:\Program Files\Kubernetes\Minikube>
C:\Program Files\Kubernetes\Minikube>eksctl get cluster --region ap-southeast-1
NAME          REGION
far-lwcluster ap-southeast-1

C:\Program Files\Kubernetes\Minikube>aws eks --region ap-southeast-1 update-kubeconfig --name far-lwcluster
Added new context arn:aws:eks:ap-southeast-1:819612740240:cluster/far-lwcluster to C:\Users\arvind.ramugade/.kube/config

C:\Program Files\Kubernetes\Minikube>
```



```

C:\Program Files\Kubernetes\Minikube>kubect1 config view
apiVersion: v1
clusters:
- cluster:
    certificate-authority-data: DATA+OMITTED
    server: https://B281A5F495624F50E70A2F5B80BBB5F5.yl4.ap-south-1.eks.amazonaws.com
    name: arn:aws:eks:ap-south-1:819612740240:cluster/lwcluster
- cluster:
    certificate-authority-data: DATA+OMITTED
    server: https://F1D498ED792166C8A8DC0775B26535A2.sk1.ap-southeast-1.eks.amazonaws.com
    name: arn:aws:eks:ap-southeast-1:819612740240:cluster/far-lwcluster
- cluster:
    certificate-authority: C:\Users\arvind.ramugade\minikube\ca.crt
    server: https://192.168.99.100:8443
    name: minikube
contexts:
- context:
    cluster: arn:aws:eks:ap-south-1:819612740240:cluster/lwcluster
    user: arn:aws:eks:ap-south-1:819612740240:cluster/lwcluster
    name: arn:aws:eks:ap-south-1:819612740240:cluster/lwcluster
- context:
    cluster: arn:aws:eks:ap-southeast-1:819612740240:cluster/far-lwcluster
    user: arn:aws:eks:ap-southeast-1:819612740240:cluster/far-lwcluster
    name: arn:aws:eks:ap-southeast-1:819612740240:cluster/far-lwcluster
- context:
    cluster: minikube
    user: minikube
    name: minikube
current-context: arn:aws:eks:ap-southeast-1:819612740240:cluster/far-lwcluster
kind: Config
preferences: {}
users:
- name: arn:aws:eks:ap-south-1:819612740240:cluster/lwcluster
  user:
    exec:
      apiVersion: client.authentication.k8s.io/v1alpha1
      args:
        - --region
        - ap-south-1
        - eks
        - get-token
        - --cluster-name
        - lwcluster
      command: aws
      env: null
- name: arn:aws:eks:ap-southeast-1:819612740240:cluster/far-lwcluster
  user:
    exec:
      apiVersion: client.authentication.k8s.io/v1alpha1
      args:
        - --region
        - ap-southeast-1
        - eks
        - get-token
        - --cluster-name
        - far-lwcluster
      command: aws
      env: null
- name: minikube
  user:
    client-certificate: C:\Users\arvind.ramugade\minikube\profiles\minikube\client.crt
    client-key: C:\Users\arvind.ramugade\minikube\profiles\minikube\client.key

```

```

C:\Program Files\Kubernetes\Minikube>kubect1 get nodes
NAME                                STATUS    ROLES    AGE   VERSION
fargate-ip-192-168-118-123.ap-southeast-1.compute.internal Ready    <none>   11m   v1.16.8-eks-e16311
fargate-ip-192-168-189-110.ap-southeast-1.compute.internal Ready    <none>   11m   v1.16.8-eks-e16311

```

To delete Fargate cluster use below command.

```
C:\Program Files\Kubernetes\Minikube>eksctl delete cluster -f D:\EKS\D2\fccluster.yml
[+] eksctl version 0.21.0
[+] using region ap-southeast-1
[+] deleting EKS cluster "far-lwcluster"
[+] deleting Fargate profile "fargate-default"
[+] deleted Fargate profile "fargate-default"
[+] deleted 1 Fargate profile(s)
[+] cleaning up LoadBalancer services
[+] 1 task: { delete cluster control plane "far-lwcluster" [async] }
[+] will delete stack "eksctl-far-lwcluster-cluster"
[+] all cluster resources were deleted
```

---